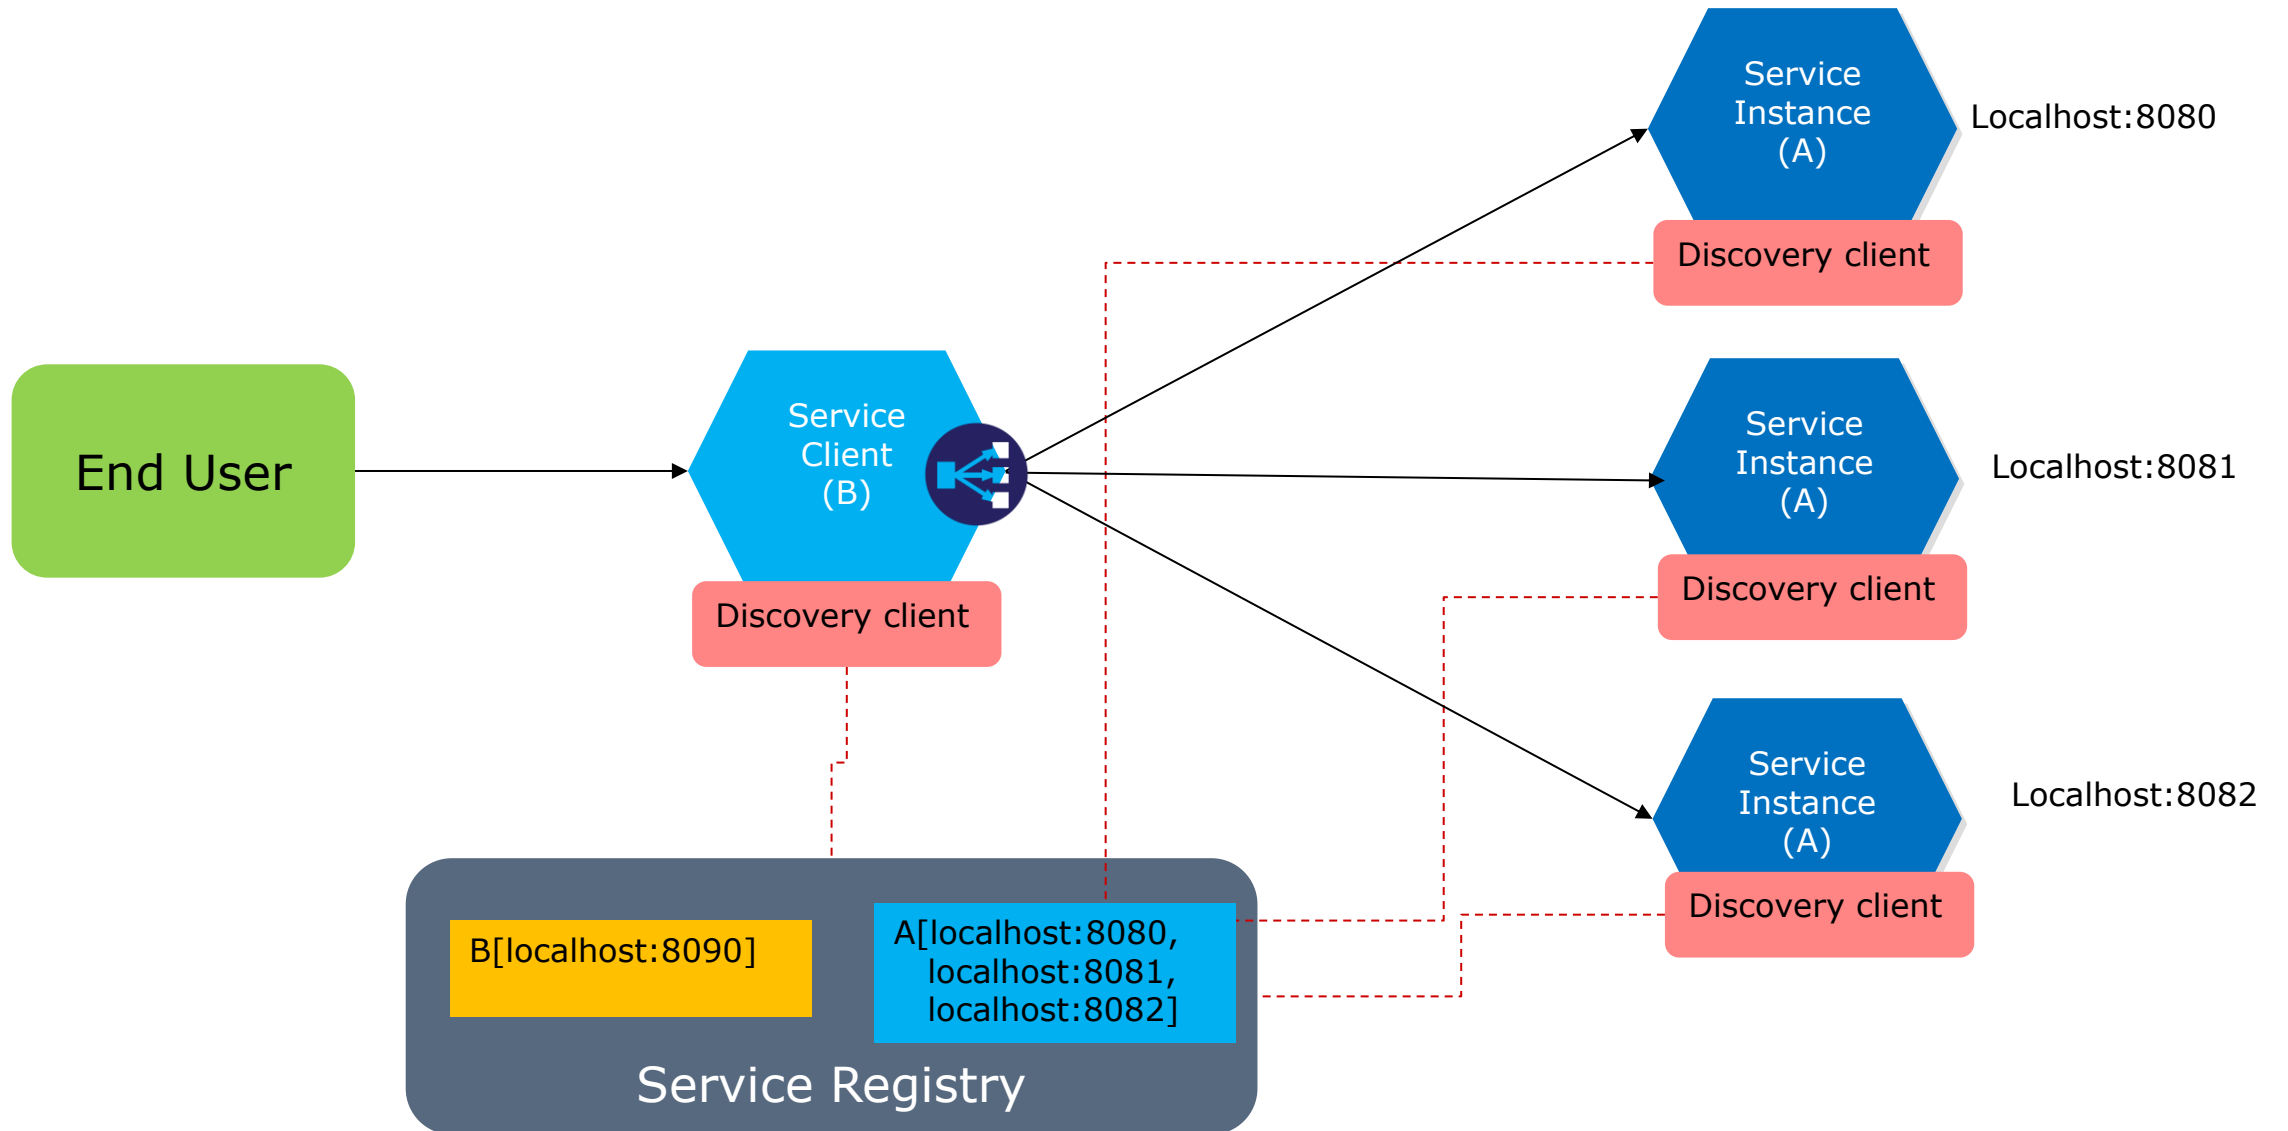


# Service Discovery and Load Balancing



# Client Side Load Balancing

---

Spring Cloud Loadbalancer

# Introduction to Client-Side Load Balancing

---

## □ **Definition:**

Client-side load balancing distributes network or application traffic across multiple servers.

## □ **Use case:**

Ensures even distribution of requests among available services in a microservice architecture.

## □ **Benefits:**

Better resource utilization, fault tolerance, and improved scalability.

---

# Client-Side Load Balancing vs. Server-Side Load Balancing

---

## ❑ **Server-Side Load Balancing:**

- Managed by the server; uses external load balancers like NGINX or HAProxy.

## ❑ **Client-Side Load Balancing:**

- Each client is aware of the available instances and chooses which one to call.

## ❑ **Example of frameworks:**

- Ribbon (deprecated).
  - Spring Cloud LoadBalancer.
-

# Spring Cloud LoadBalancer Overview

---

- ❑ Replaced Netflix Ribbon in Spring Cloud.
  - ❑ Integrated with Spring Boot and Spring Cloud microservices.
  - ❑ Automatically distributes requests among instances of a service.
-

# How Spring Cloud LoadBalancer Works

---

- ❑ The client retrieves a list of service instances from a **Service Discovery** mechanism (like Eureka).
  - ❑ Balancing logic runs on the client to decide which instance to invoke.
  - ❑ LoadBalancerClient manages the request distribution.
-

# Configuration of Spring Cloud LoadBalancer

---

- ❑ Add Spring Cloud LoadBalancer dependency:
  - ❑ Ensure Service Discovery integration (Eureka, Consul, etc.).
  - ❑ Enable client-side load balancing by annotating the RestTemplate or WebClient.
-

# Client-Side Load Balancing in Spring Boot

---

## □ Example using RestTemplate:

```
@Bean
@LoadBalanced
public RestTemplate restTemplate() {
    return new RestTemplate();
}
```

## □ Example using WebClient:

```
@Bean
@LoadBalanced
public WebClient.Builder webClientBuilder() {
    return WebClient.builder();
}
```



# Service Discovery Integration

---

- Common integrations:
  - **Eureka:** Popular for service registry and discovery.
  - **Consul:** A more robust alternative for service discovery.
- The LoadBalancer gets available instances from these registries.

# Load Balancing Strategies

---

- ❑ **Round Robin:** Distributes requests evenly in a circular order.
  - ❑ **Random:** Selects an instance at random.
  - ❑ **Weighted:** Requests are distributed based on pre-defined weights of instances.
-

# Customizing Load Balancing Behavior

---

- Custom rule implementation for specific needs:

@Bean

```
public ReactorLoadBalancer<ServiceInstance> customLoadBalancer(Environment environment) {  
    String name = environment.getProperty(LoadBalancerClientFactory.PROPERTY_NAME);  
    return new RandomLoadBalancer(name);  
}
```

---

# Monitoring & Troubleshooting

---

- Use tools like Spring Boot Actuator to monitor service instances.
  - Common issues:
    - Misconfiguration of Service Discovery.
    - Inconsistent request distribution.
    - Health-check failures of instances.
-

# Best Practices

---

- ❑ Use retry mechanisms to handle failed requests.
  - ❑ Combine client-side load balancing with circuit breakers (e.g., Resilience4j).
  - ❑ Ensure health checks are configured for all service instances.
-

# Service Discovery

---

Eureka Server

# The Need for Service Discovery

---

- ❑ Problem: As microservices scale, services need to find and communicate with each other dynamically.
  - ❑ Manual service registration vs. automated service discovery
  - ❑ Why hard-coding service locations doesn't scale
-

# What is Service Discovery?

---

- ❑ Definition: A mechanism where services register themselves with a central registry and clients discover the service dynamically.
  - ❑ Importance in Microservices architecture
  - ❑ Types of Service Discovery: Client-Side and Server-Side
-



# Client-Side Service Discovery

---

- ❑ Definition and example
  - ❑ How clients (microservices) look up services directly from the registry
  - ❑ Diagram: Service Client interacting directly with Service Registry
  - ❑ Tools: Netflix Eureka, Consul, Zookeeper
-

# Server-Side Service Discovery

---

- Definition and example
  - Service registry is used by a load balancer or gateway
  - Diagram: Clients interacting with a load balancer that queries the registry
  - Tools: Kubernetes (Kube-DNS), AWS ELB, Nginx with Consul
-

# Introducing Netflix Eureka

---

- What is Netflix Eureka?
  - Role in the Spring ecosystem
  - Features: Service Registration, Heartbeats, Service Discovery
-

# Setting up Eureka Server in Spring Boot

---

- ❑ Add dependencies: `spring-cloud-starter-netflix-eureka-server`
- ❑ Configurations for Eureka server:

server:

port: 8761

eureka:

client:

register-with-eureka: false

fetch-registry: false

---

# Setting up Eureka Client in Spring Boot

---

- ❑ Add dependencies: `spring-cloud-starter-netflix-eureka-client`
- ❑ Service registration configuration:

**eureka:**

**client:**

**service-url:**

**defaultZone:** `http://localhost:8761/eureka/`

- ❑ Registering multiple services in Eureka
-

# Service Discovery in Action

---

- How services register themselves on the Eureka Server
  - How clients discover other services via Eureka
  - Live demo: Register two microservices and use Eureka to discover them
-

# Fault Tolerance in Service Discovery

---

- ❑ Circuit breaker pattern with Spring Cloud and Resilience4j
  - ❑ Handling service unavailability
  - ❑ Fall-back mechanisms
-

# Server-Side Discovery with Spring Cloud Gateway

---

- ❑ Integrating Eureka with Spring Cloud Gateway
  - ❑ Dynamic routing based on service registry
  - ❑ Config example for using Spring Cloud Gateway with Eureka
-



# Scaling and Performance Considerations

---

- Horizontal scaling of Eureka servers
  - High availability with multiple Eureka instances
  - Performance tuning in large-scale applications
-

# Alternatives to Eureka

---

- Consul
  - Zookeeper
  - Kubernetes Service Discovery
-

# Conclusion

---

- Recap of Service Discovery benefits
  - Overview of Eureka in Spring Boot
  - Final thoughts on best practices
-

# Circuit Breaker Pattern

---

Resilience4j in Spring Boot

Enhancing Microservice Resilience

V-1.0

# Agenda

---

- ❑ What is the Circuit Breaker Pattern?
  - ❑ Resilience in Microservices Architecture
  - ❑ Introduction to Resilience4j
  - ❑ Resilience4j Circuit Breaker with Spring Boot
  - ❑ Demonstration: Implementing Circuit Breaker
  - ❑ Key Takeaways
-

# What is the Circuit Breaker Pattern?

---

- **Definition:** A design pattern that helps prevent cascading failures in distributed systems by temporarily stopping the execution of requests to an external service that's failing.
  - **Phases:**
    - **Closed:** Requests flow normally.
    - **Open:** Requests are blocked for a time.
    - **Half-Open:** Limited requests are sent to check if the service has recovered.
-

# Why Use Circuit Breakers?

---

- ❑ Prevent service overloads.
  - ❑ Improve application resilience.
  - ❑ Enhance fault tolerance in distributed architectures.
  - ❑ Handle failures gracefully.
-

# Resilience in Microservices

---

- ❑ Challenges in microservice communication:
  - **Latency issues**
  - **Service failures**
  - **Network partitions**
- ❑ Solutions: Circuit Breakers, Retries, Bulkheads, Rate Limiters



# Introduction to Resilience4j

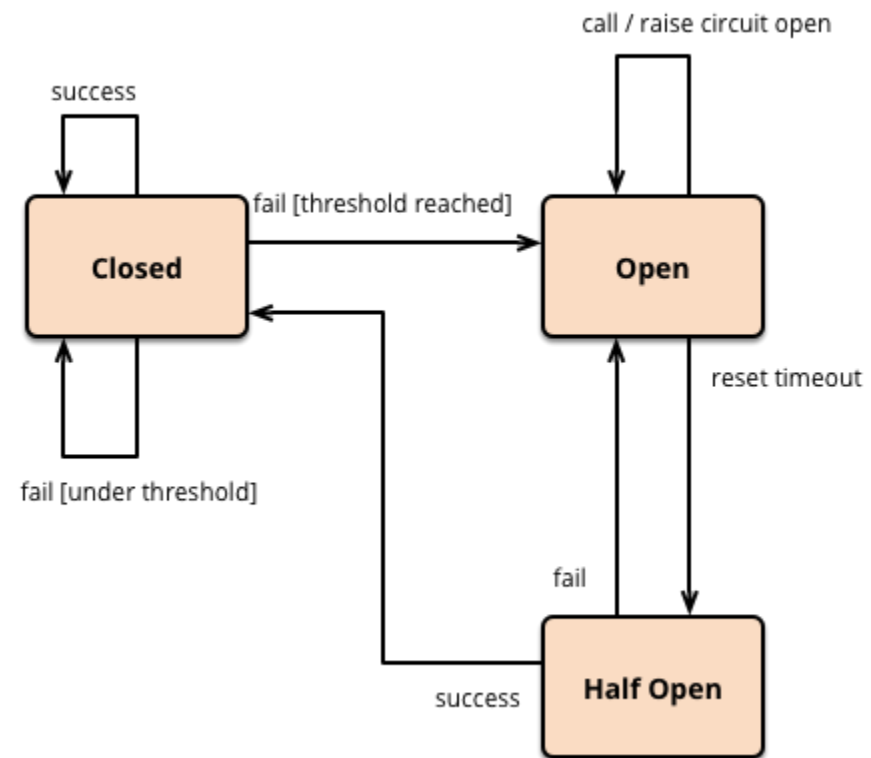
---

- ❑ **Resilience4j**: A lightweight, easy-to-use fault tolerance library for Java 8+ that provides implementations for circuit breakers, rate limiters, retries, etc.
  - ❑ Modular, compared to Netflix Hystrix (retired).
  - ❑ Key modules:
    - CircuitBreaker
    - Retry
    - RateLimiter
    - Bulkhead
    - TimeLimiter
-

# How Circuit Breaker Works in Resilience4j

---

- Circuit Breaker transitions between:
    - **Closed**: All requests pass through.
    - **Open**: Requests are blocked for a time.
    - **Half-Open**: Test requests pass through to see if the service has recovered.
  - Configurable parameters:
    - Failure rate threshold
    - Wait duration before opening
    - Sliding window size
-



# Integrating Resilience4j with Spring Boot

---

- ❑ Resilience4j provides Spring Boot Starter libraries.
  - ❑ Configuration via application.properties or YAML files.
  - ❑ Integration steps:
    - ❑ Add the Resilience4j dependency.
    - ❑ Configure circuit breaker properties.
    - ❑ Annotate methods with @CircuitBreaker.
-

# Dependency Management


---

```
<dependency>  
    <groupId>io.github.resilience4j</groupId>  
    <artifactId>resilience4j-spring-boot2</artifactId>  
    <version>1.7.0</version>  
</dependency>
```

---

# Properties

```
resilience4j.circuitbreaker:  
  instances:  
    message-service:  
      registerHealthIndicator: true  
      slidingWindowSize: 10  
      permittedNumberOfCallsInHalfOpenState: 3  
      slidingWindowType: COUNT_BASED  
      minimumNumberOfCalls: 5  
      waitDurationInOpenState: 5s  
      failureRateThreshold: 33.3  
      automaticTransitionFromOpenToHalfOpenEnabled: true
```



# Code Example

---

```
@CircuitBreaker(name = "myService", fallbackMethod = "fallbackMethod")
public String myService() {
    // Call to an external service
}

public String fallbackMethod(Throwable t) {
    return "Service is currently unavailable.";
}
```

# Demo: Circuit Breaker Implementation in Spring Boot

---

- Short demo of how to create a circuit breaker-enabled service in Spring Boot.
  - **Steps:**
    - Setup dependencies.
    - Create a service class.
    - Configure circuit breaker.
    - Test with an external service failure.
-



# Circuit Breaker Metrics and Monitoring

---

- ❑ Micrometer integration: Monitor circuit breaker states.
- ❑ Actuator: Use Spring Boot Actuator for health checks and exposing circuit breaker metrics.
- ❑ Example:

`/actuator/metrics/resilience4j.circuitbreaker.state`

---

# Fallback Strategies

---

- ❑ Fallback methods ensure the application doesn't crash.
  - ❑ Fallbacks can return default responses or cached data.
  - ❑ Example:
    - Static response
    - Another service call
-

# Key Takeaways

---

- ❑ Circuit breaker pattern improves the resilience of microservices.
  - ❑ Resilience4j offers a modular and efficient way to implement circuit breakers.
  - ❑ Easy integration with Spring Boot via annotations and configuration.
-

---

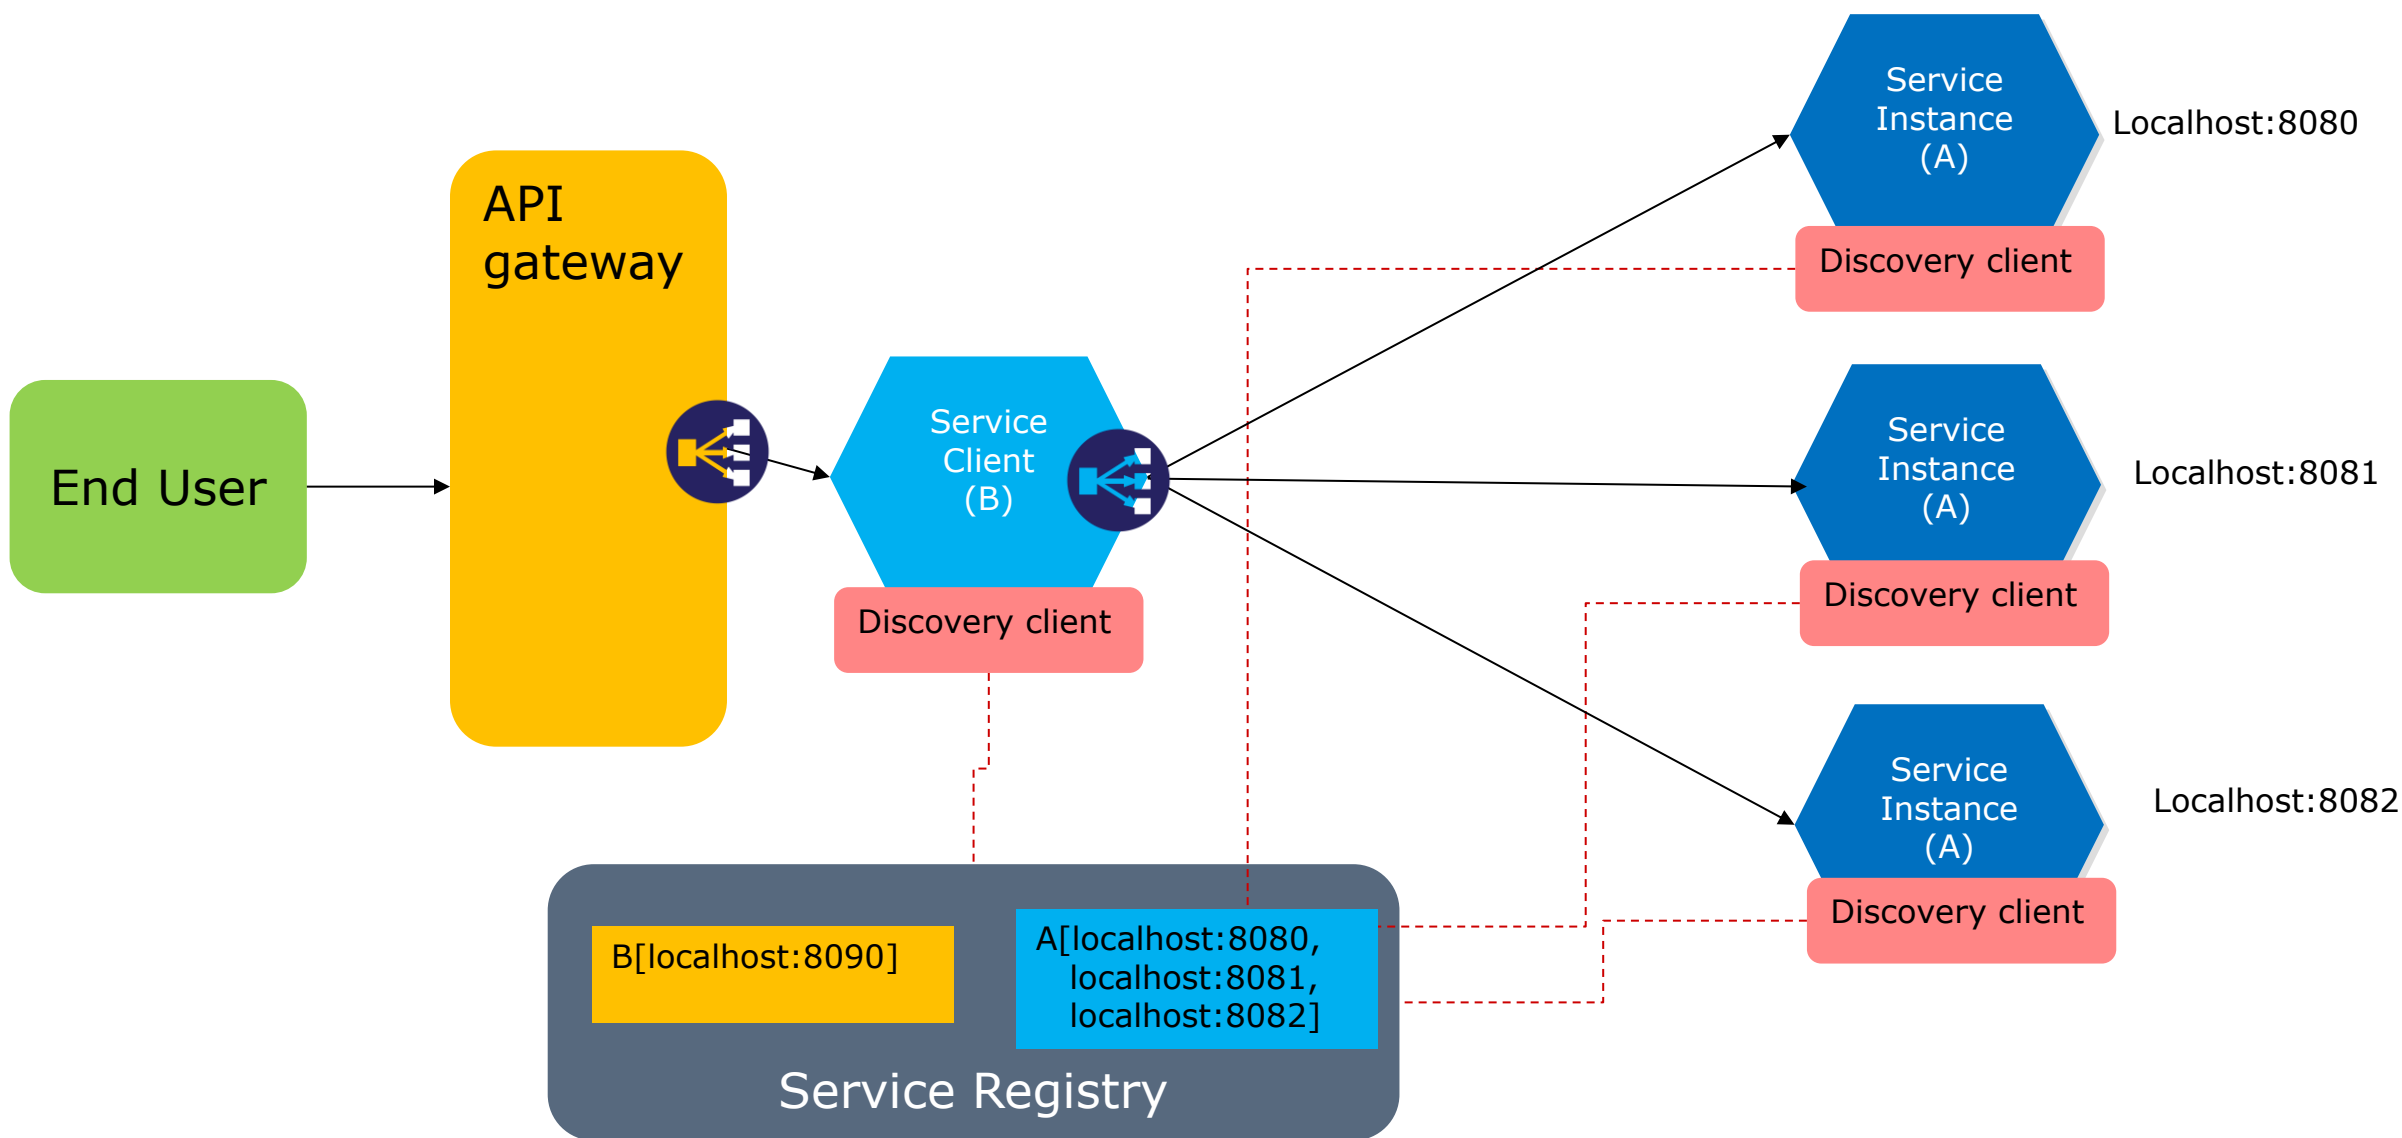
□ Questions?

---

# API Gateway

---

# Service Discovery and Load Balancing



# Introduction to API Gateway Pattern

---

## □ What is an API Gateway?

- A server that acts as an entry point for all client requests.
- It routes requests to the appropriate backend microservices.

## □ Why use an API Gateway?

- Centralized access control
  - Cross-cutting concerns (logging, security, caching)
  - Reduces complexity for clients
  - Helps decouple front-end clients from multiple microservices
-

# Key Features of an API Gateway

---

## ❑ Request Routing

- Directs client requests to the right microservice

## ❑ Load Balancing

- Distributes incoming traffic across microservices

## ❑ Security

- Authentication, authorization, and SSL termination

## ❑ Rate Limiting

- Controls the number of requests per client

## ❑ Logging & Monitoring

- Tracks request flows, error logs, and metrics
-



# Spring Cloud Gateway Overview

---

## □ What is Spring Cloud Gateway?

- A project built on top of Spring Boot for building API gateways.
- Provides features like routing, filters, security, and more.

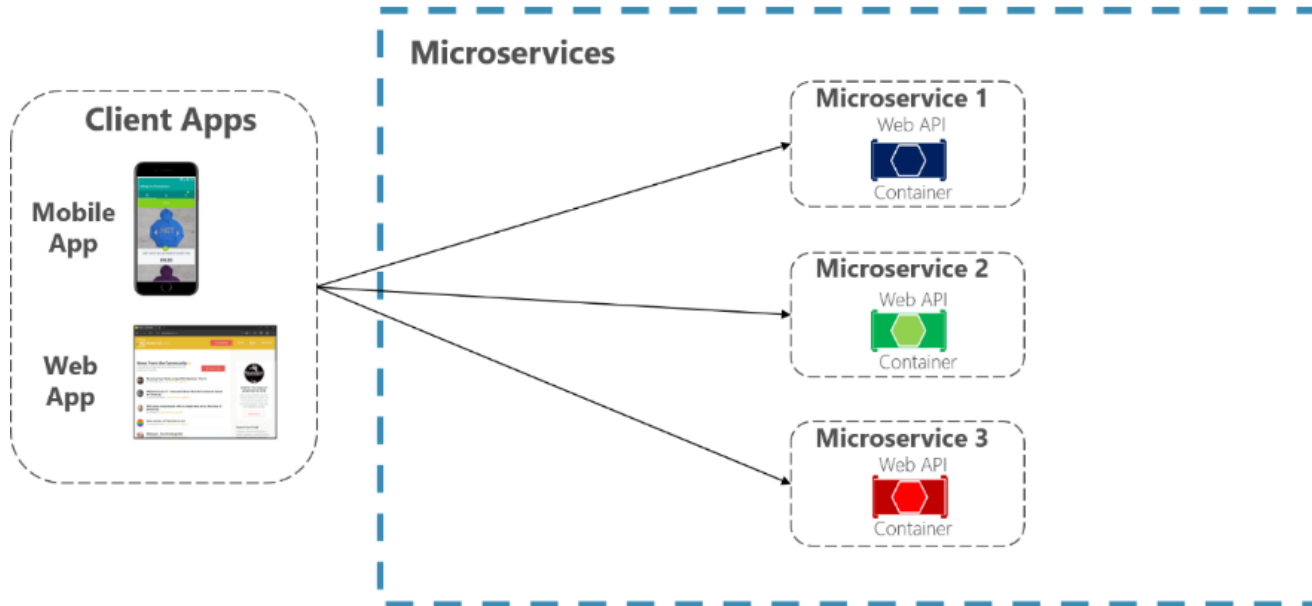
## □ Why use Spring Cloud Gateway?

- Built on top of Project Reactor for reactive programming
  - Easy integration with Spring ecosystem
-

# Direct Access: No API Gateway

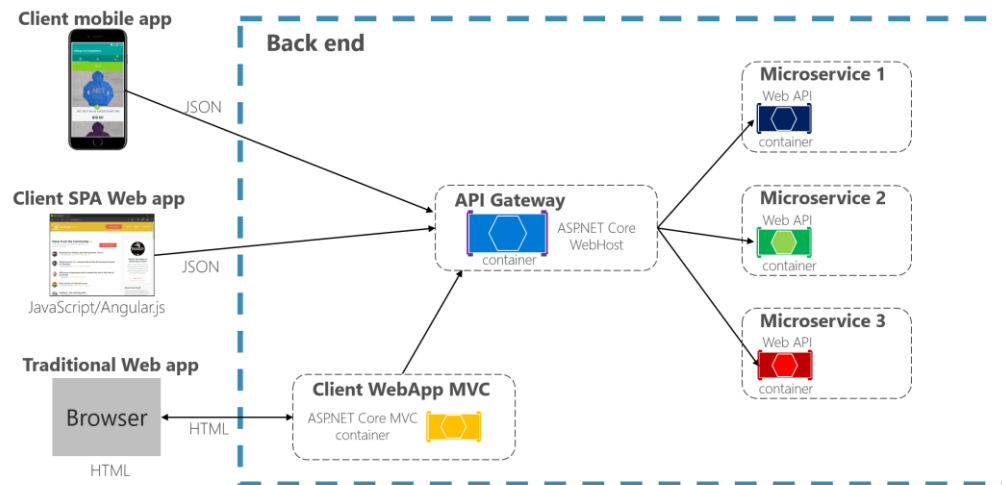
---

## Direct Client-To-Microservice communication Architecture



# Architecture Diagram

Using a single custom **API Gateway service**



Using multiple **API Gateways / BFF**

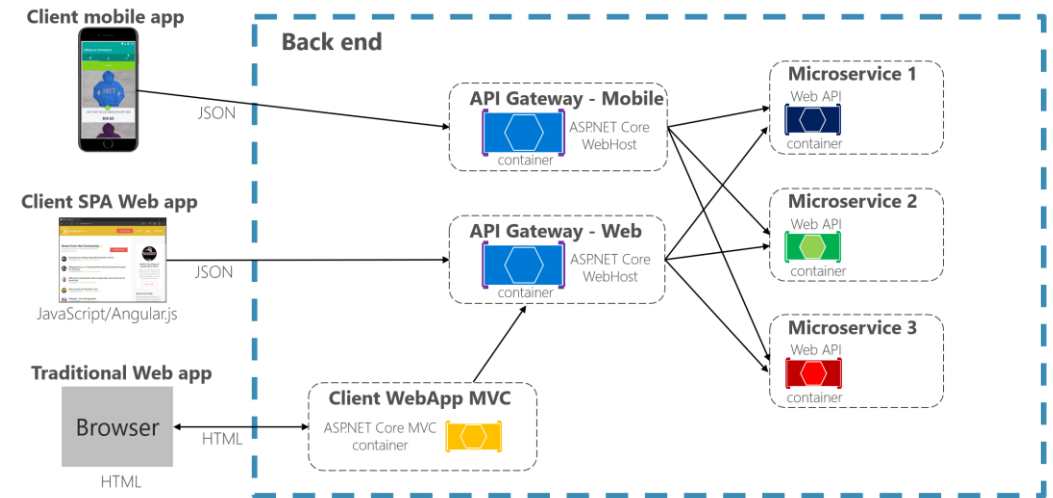


Image Courtesy: Microsoft

---

## □ **API Gateway Flow in Spring Boot**

### **1. Client Request**

### **2. API Gateway (Spring Cloud Gateway)**

1. Request routed to appropriate microservice

### **3. Microservice Communication**

1. Response sent back through the gateway to the client

---

- 
- ❑ Setting up Spring Cloud Gateway
  - ❑ Dependencies in pom.xml:
  - ❑ xml
  - ❑ Copy code
  - ❑ <dependency>
  - ❑     <groupId>org.springframework.cloud</groupId>
  - ❑     <artifactId>spring-cloud-starter-gateway</artifactId>
  - ❑ </dependency>
  - ❑ Main Class Setup:
  - ❑ java
  - ❑ Copy code
  - ❑ @SpringBootApplication
  - ❑ public class ApiGatewayApplication {
  - ❑     public static void main(String[] args) {
  - ❑         SpringApplication.run(ApiGatewayApplication.class, args);
  - ❑     }
  - ❑ }
-

- 
- ❑ Configuring Routes in API Gateway
  - ❑ Example of Route Configuration in application.yml:
  - ❑ yml
  - ❑ Copy code
  - ❑ spring:
  - ❑   cloud:
  - ❑     gateway:
  - ❑       routes:
  - ❑         - id: user-service
  - ❑         uri: http://localhost:8081
  - ❑         predicates:
  - ❑         - Path=/users/\*\*
  - ❑         - id: order-service
  - ❑         uri: http://localhost:8082
  - ❑         predicates:
  - ❑         - Path=/orders/\*\*
  - ❑ Explanation:
  - ❑ Define route IDs, URIs, and paths for microservices.
-

- 
- ❑ Filters in API Gateway
  - ❑ Global Filters:
  - ❑ Applied to all routes.
  - ❑ Examples: Security, logging.
  - ❑ Route-Specific Filters:
  - ❑ Applied to specific routes.
  - ❑ Example of a Filter:
  - ❑ yaml
  - ❑ Copy code
  - ❑ - id: order-service
  - ❑ uri: http://localhost:8082
  - ❑ filters:
  - ❑ - AddRequestHeader=Order,Service
-

- 
- ❑ Load Balancing with API Gateway
  - ❑ Enable Load Balancing:
  - ❑ Use Spring Cloud's spring-cloud-starter-loadbalancer.
  - ❑ Example Configuration:
  - ❑ yaml
  - ❑ Copy code
  - ❑ spring:
  - ❑   cloud:
  - ❑     gateway:
  - ❑       discovery:
  - ❑         locator:
  - ❑         enabled: true
-



---

## ❑ **Security in API Gateway**

- **Authentication & Authorization:**

- Use Spring Security with OAuth2 or JWT.

- **Example Configuration:**

- Add JWT validation to requests before forwarding them to microservices.
-

- 
- ❑ Rate Limiting in API Gateway
  - ❑ Example Configuration for Rate Limiting:
  - ❑ yaml
  - ❑ Copy code
  - ❑ - id: rate\_limiter
  - ❑ uri: http://localhost:8080
  - ❑ filters:
  - ❑ - name: RequestRateLimiter
  - ❑ args:
  - ❑ redis-rate-limiter:
  - ❑ replenishRate: 10
  - ❑ burstCapacity: 20
-

---

## □ **API Gateway Best Practices**

- **Handle timeouts and retries effectively.**
  - **Ensure security policies are applied centrally.**
  - **Monitor the gateway for performance bottlenecks.**
  - **Implement caching for faster responses.**
-

---

## □ Conclusion

- **Recap:**
  - API Gateway pattern simplifies microservice management.
  - Spring Cloud Gateway provides a powerful tool for implementing it in Spring Boot.